

# Actividad modulo 27 - Regresión lineal multivariable

## Descripción de los campos

- Price (#)
  - Area (#)
  - Bedroom (#)
  - Bathroom (#)
  - Stories (#)
  - Mainroad (yes/no)
  - guestroom (yes/no)
  - basement (yes/no)
  - hotwaterheating (yes/no)
  - airconditioning (yes/no)
  - parking (#)
  - prefarea(yes/no)
  - furnishingstatus (furnished, semi-furnished, unfurnished)
- En este caso la idea seria predecir el precio de las casas, basado en los 545 registros del dataset

```
In [ ]: import pandas as pd
import numpy as np
import os
import seaborn as sns
import matplotlib.pyplot as plt

os.chdir('E:\WORK IN PROGRESS\Data Analytics course\parte 2 python\week 27')

df = pd.read_csv('Housing.csv')
```

## Exploratory Data Analysis (EDA)

```
In [ ]: df.sample(5)
```

Out [ ]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
153	5530000	3300	3	3	2	yes	no	yes	no	no	0	no	semi-furnished
256	4480000	4000	3	1	2	yes	no	no	no	no	2	no	furnished
507	2590000	3600	2	1	1	yes	no	no	no	no	0	no	unfurnished
268	4382000	4950	4	1	2	yes	no	no	no	yes	0	no	semi-furnished
227	4690000	6000	2	1	1	yes	no	yes	no	yes	1	no	furnished

```
In [ ]: # Validar cuantas celdas en total son nulas
print ('Existen', df.isnull().sum().sum(), 'valores no existentes')
```

Existen 0 valores no existentes

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   price                545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms             545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom            545 non-null   object
7   basement             545 non-null   object
8   hotwaterheating      545 non-null   object
9   airconditioning      545 non-null   object
10  parking              545 non-null   int64
11  prefarea             545 non-null   object
12  furnishingstatus     545 non-null   object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

```
In [ ]: pd.set_option('display.float_format', lambda x: '%.3f' % x)
```

```
In [ ]: # Obtener los principales estadísticos por cada uno
df.describe().T
```

Out [ ]:

	count	mean	std	min	25%	50%	75%	max
price	545.000	4766729.248	1870439.616	1750000.000	3430000.000	4340000.000	5740000.000	13300000.000
area	545.000	5150.541	2170.141	1650.000	3600.000	4600.000	6360.000	16200.000
bedrooms	545.000	2.965	0.738	1.000	2.000	3.000	3.000	6.000
bathrooms	545.000	1.286	0.502	1.000	1.000	1.000	2.000	4.000
stories	545.000	1.806	0.867	1.000	1.000	2.000	2.000	4.000
parking	545.000	0.694	0.862	0.000	0.000	0.000	1.000	3.000

```
In [ ]: # Se ve el total de los valores únicos por cada columna
# Al parecer es un buen data set de datos, no tiene tanta variabilidad en las respuestas, lo que hace que se a mas facil su prediccion
# Además se confirma en los valores yes/no y furnished, que tienen 2 y 3 valores, correspondientemente
df.nunique()
```

Out [ ]:

```
price      219
area       284
bedrooms    6
bathrooms   4
stories     4
mainroad    2
guestroom   2
basement    2
hotwaterheating  2
airconditioning  2
parking     4
prefarea    2
furnishingstatus  3
dtype: int64
```

## Feature Engineering

```
In [ ]: # Antes de seguir se cambian las columnas que tienen valores en caneda a numeros
# yes/no por 1/0
# furnished/semi-furnished/unfurnished:1/2/3
object_columns=[col for col in df.columns if (df[col].dtypes=='object')]
object_columns.remove('furnishingstatus')
for j in object_columns:
    df[j]=df[j].map({'yes':1,'no':0})
df['furnishingstatus']=df['furnishingstatus'].map({'furnished':1,'semi-furnished':2,'unfurnished':3})
```

```
In [ ]: # Se revisa la nueva estructura
df.sample(5)
```

Out [ ]:

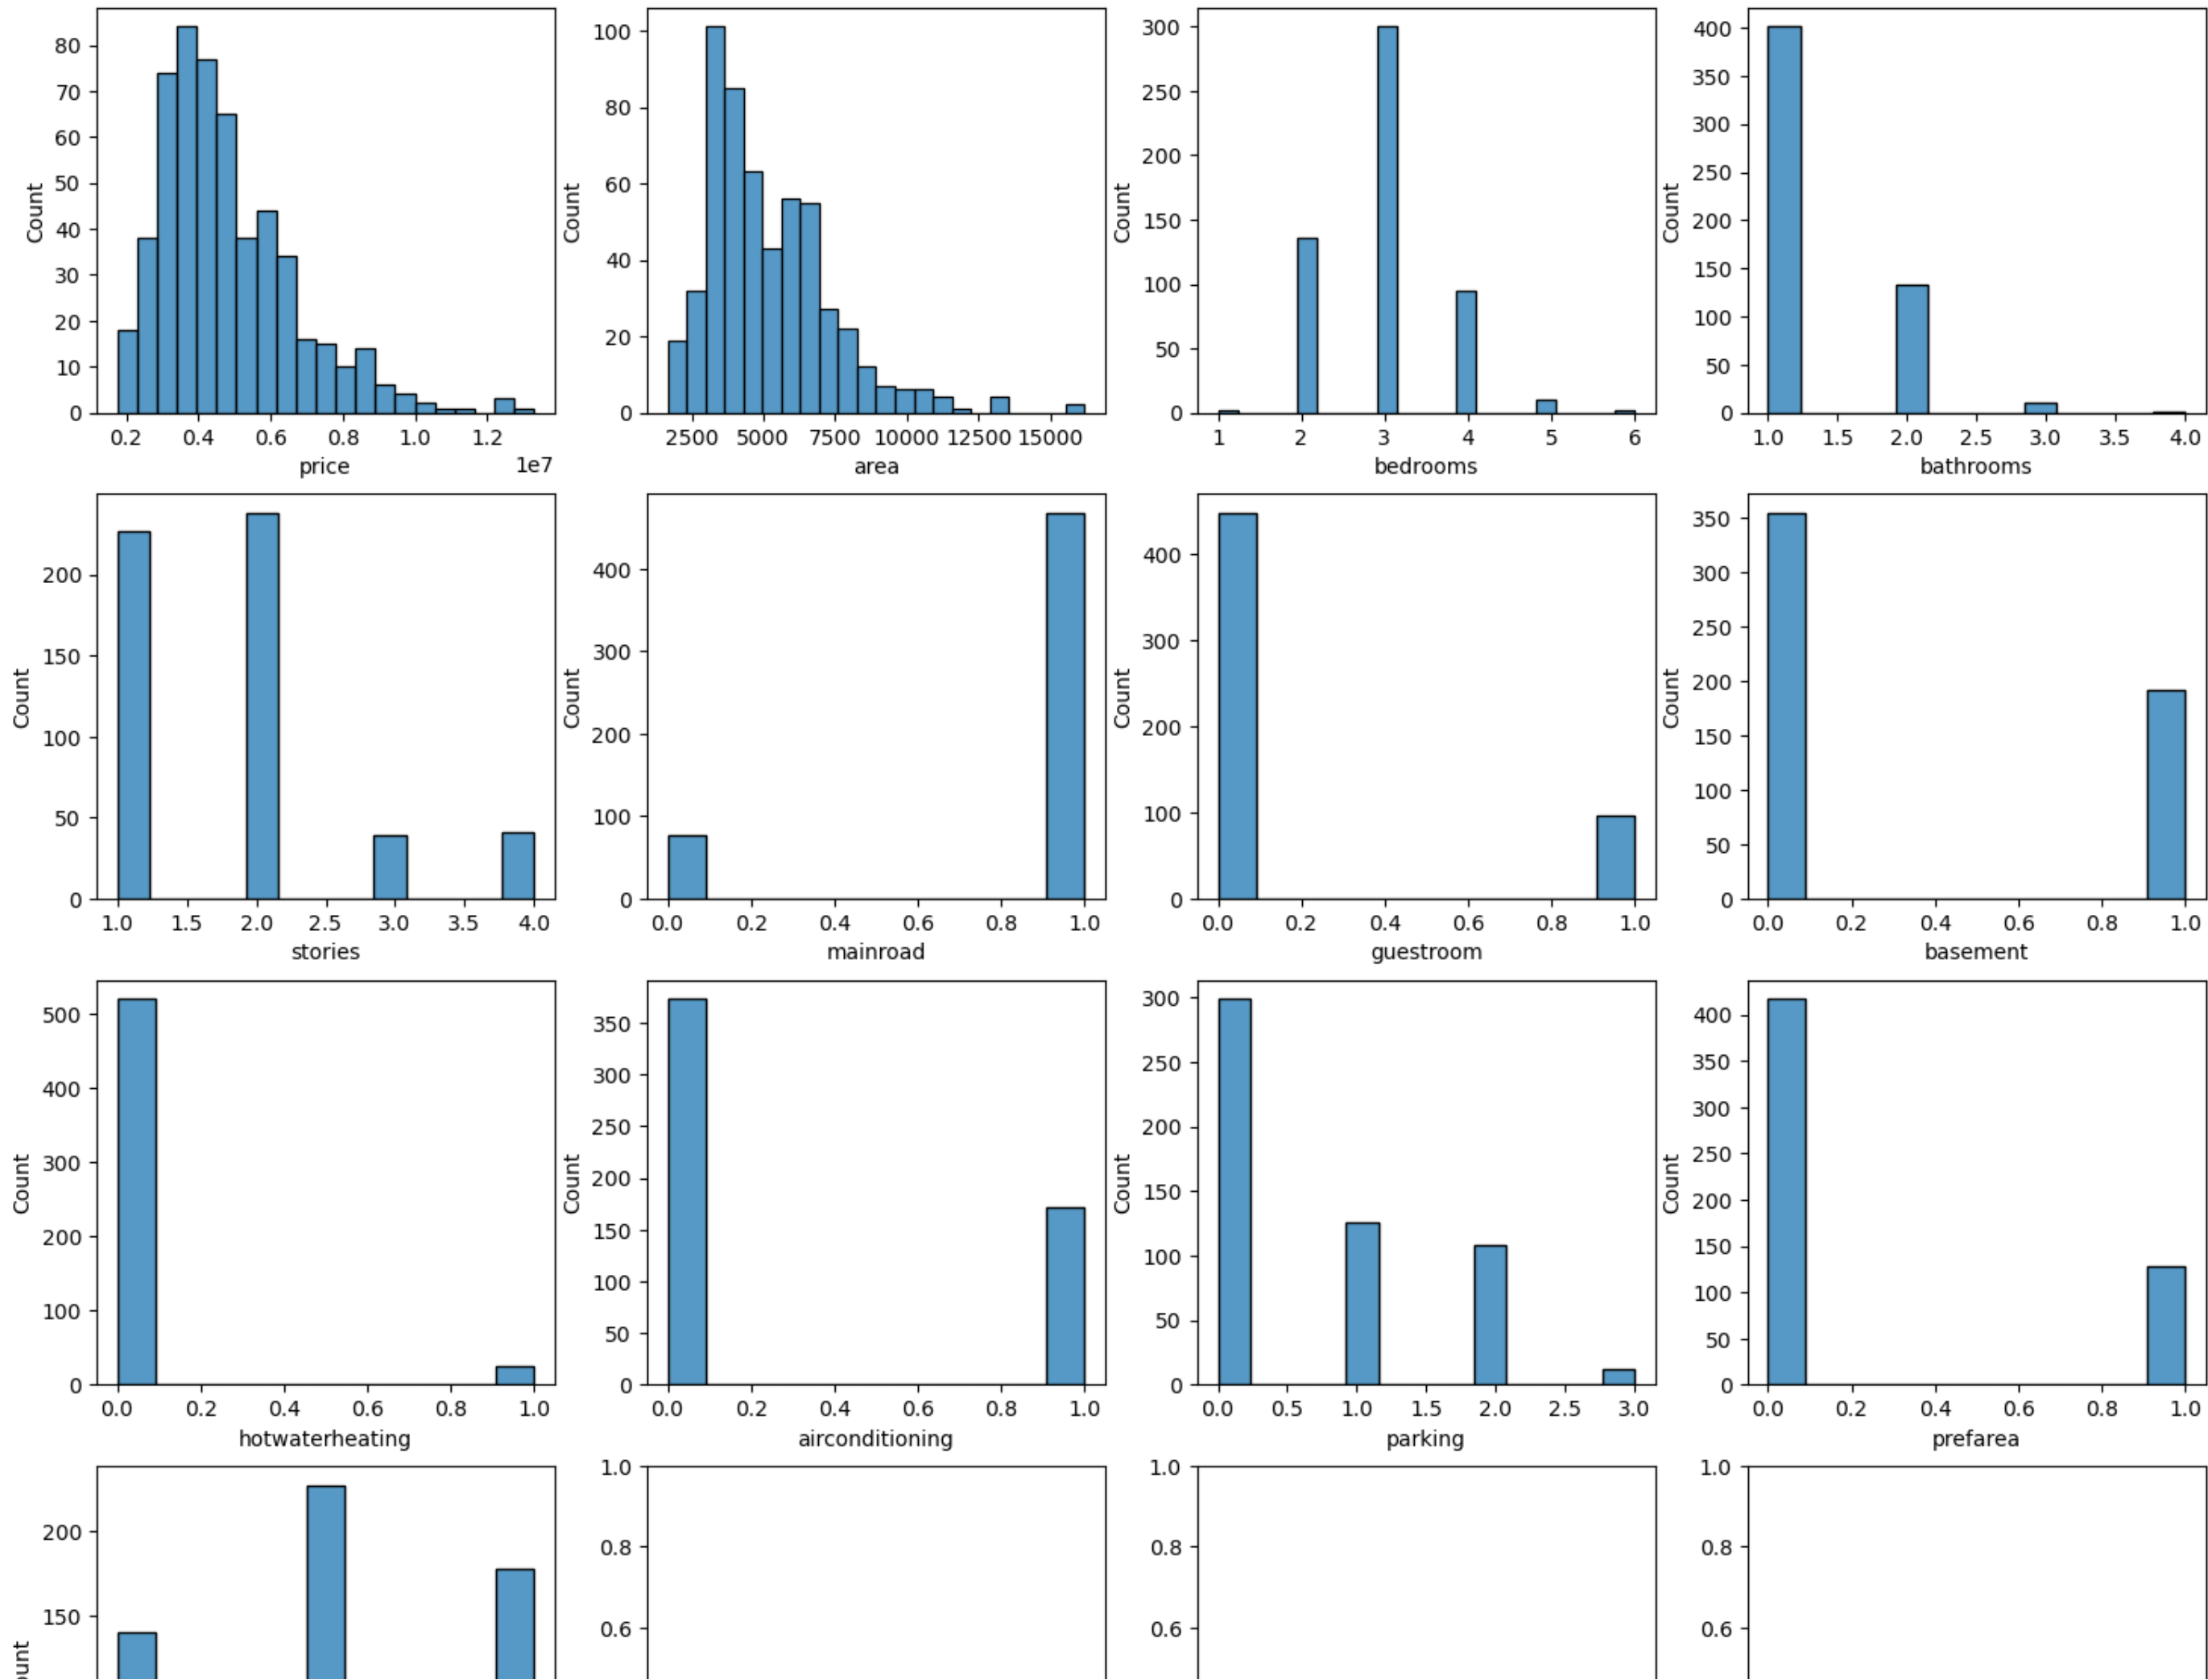
	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
34	8120000	6840	5	1	2	1	1	1	0	1	1	0	1
108	6107500	3240	4	1	3	1	0	0	0	0	1	0	2
290	4200000	2610	4	3	2	0	0	0	0	0	0	0	2
510	2520000	2880	3	1	1	0	0	0	0	0	0	0	3
483	2940000	6615	3	1	2	1	0	0	0	0	0	0	2

```
In [ ]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   price               545 non-null   int64
1   area                545 non-null   int64
2   bedrooms            545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   int64
6   guestroom            545 non-null   int64
7   basement             545 non-null   int64
8   hotwaterheating      545 non-null   int64
9   airconditioning      545 non-null   int64
10  parking              545 non-null   int64
11  prefarea             545 non-null   int64
12  furnishingstatus     545 non-null   int64
dtypes: int64(13)
memory usage: 55.5 KB
```


```
In [ ]: # Grafico exploratorio de todas las columnas
cols_num=df.columns.to_list()

fig, axes = plt.subplots(nrows=4,ncols=4, figsize=(18,16))
for i, column in enumerate(cols_num):
    sns.histplot(df[column],ax=axes[i//4,i%4],kde=False)
```



Analisis Correlacion / Bivariado

```
In [ ]: # Grafica de Correlacion
plt.figure(figsize=(22,8))
corr_df = corr= df.corr(method='pearson')
df_lt=corr_df.where(np.tril(np.ones(corr_df.shape)).astype(bool))
hmap=sns.heatmap(df_lt, cmap='YlGnBu',annot=True)
```



Insights

- No se encuentra una correlación significativa entre las variables independientes. Por lo cual no es necesario eliminar ninguna de ellas.

## Balance de clases

```
In [ ]: columns= df.columns.to_list()
categ_columns = [column for column in columns if (column != 'area') & (column !='price')]
categ_columns
```

```
Out[ ]: ['bedrooms',
'bathrooms',
'stores',
'mainroad',
'guestroom',
'basement',
'hotwaterheating',
'airconditioning',
'parking',
'prefarea',
'furnishingstatus']
```

```
In [ ]: for column in categ_columns:
    helper = df.groupby(column)[column].count().rename('total').reset_index()

    helper['percent']=(helper['total']/helper['total'].sum())*100
    print('\n' + column)
    print(helper)
```

bedrooms			
	bedrooms	total	percent
0	1	2	0.367
1	2	136	24.954
2	3	300	55.046
3	4	95	17.431
4	5	10	1.835
5	6	2	0.367
bathrooms			
	bathrooms	total	percent
0	1	401	73.578
1	2	133	24.404
2	3	10	1.835
3	4	1	0.183
stories			
	stories	total	percent
0	1	227	41.651
1	2	238	43.670
2	3	39	7.156
3	4	41	7.523
mainroad			
	mainroad	total	percent
0	0	77	14.128
1	1	468	85.872
guestroom			
	guestroom	total	percent
0	0	448	82.202
1	1	97	17.798
basement			
	basement	total	percent
0	0	354	64.954
1	1	191	35.046
hotwaterheating			
	hotwaterheating	total	percent
0	0	520	95.413
1	1	25	4.587
airconditioning			
	airconditioning	total	percent
0	0	373	68.440
1	1	172	31.560
parking			
	parking	total	percent
0	0	299	54.862
1	1	126	23.119
2	2	108	19.817
3	3	12	2.202
prefarea			
	prefarea	total	percent
0	0	417	76.514
1	1	128	23.486
furnishingstatus			
	furnishingstatus	total	percent
0	1	140	25.688

1	2	227	41.651
2	3	178	32.661

## Parking

- cambio los valores de '3' por '2' asi la opcion 2 se entiende como + de 1 parking

```
In [ ]: df['parking'].replace({3:2},inplace=True)
```

## Bathrooms

- cambio los valores de '3' y '4' por '2' asi la opcion 2 se entiende como + de 1 banos

```
In [ ]: df['bathrooms'].replace({3:2,4:2},inplace=True)
```

## Bedrooms

- cambio los valores de '5' y '6' por '4' asi la opcion 4 se entiende como + de 3 banos
- elimino los valores de '1' bedroom

```
In [ ]: df['bedrooms'].replace({5:4,6:4},inplace=True)

df.drop(df[df['bedrooms']==1].index,inplace=True)
```

```
In [ ]: for column in categ_columns:
    helper = df.groupby(column)[column].count().rename('total').reset_index()

    helper['percent']=(helper['total']/helper['total'].sum())*100
    print('\n' + column)
    print(helper)
```



bedrooms			
	bedrooms	total	percent
0	2	136	25.046
1	3	300	55.249
2	4	107	19.705

bathrooms			
	bathrooms	total	percent
0	1	399	73.481
1	2	144	26.519

stories			
	stories	total	percent
0	1	225	41.436
1	2	238	43.831
2	3	39	7.182
3	4	41	7.551

mainroad			
	mainroad	total	percent
0	0	76	13.996
1	1	467	86.004

guestroom			
	guestroom	total	percent
0	0	446	82.136
1	1	97	17.864

basement			
	basement	total	percent
0	0	352	64.825
1	1	191	35.175

hotwaterheating			
	hotwaterheating	total	percent
0	0	518	95.396
1	1	25	4.604

airconditioning			
	airconditioning	total	percent
0	0	371	68.324
1	1	172	31.676

parking			
	parking	total	percent
0	0	297	54.696
1	1	126	23.204
2	2	120	22.099

prefarea			
	prefarea	total	percent
0	0	415	76.427
1	1	128	23.573

furnishingstatus			
	furnishingstatus	total	percent
0	1	139	25.599
1	2	227	41.805
2	3	177	32.597

Generación de variables dummy en el set de datos

```
In [ ]: df.nunique()
```

```
Out[ ]: price      219
area      284
bedrooms   3
bathrooms  2
stories    4
mainroad   2
guestroom  2
basement   2
hotwaterheating  2
airconditioning  2
parking    3
prefarea   2
furnishingstatus  3
dtype: int64
```

```
In [ ]: for col in categ_columns:
        df[col]= df[col].astype('category')

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 543 entries, 0 to 544
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   price           543 non-null   int64
1   area            543 non-null   int64
2   bedrooms        543 non-null   category
3   bathrooms        543 non-null   category
4   stories          543 non-null   category
5   mainroad         543 non-null   category
6   guestroom        543 non-null   category
7   basement         543 non-null   category
8   hotwaterheating  543 non-null   category
9   airconditioning  543 non-null   category
10  parking          543 non-null   category
11  prefarea         543 non-null   category
12  furnishingstatus 543 non-null   category
dtypes: category(11), int64(2)
memory usage: 20.0 KB
```

```
In [ ]: df= pd.get_dummies(df,drop_first=True)
df.sample(5)
```

	price	area	bedrooms_3	bedrooms_4	bathrooms_2	stories_2	stories_3	stories_4	mainroad_1	guestroom_1	basement_1	hotwaterheating_1	airconditioning_1	parking_1	parking_2	prefarea_1	furnishingstatus_2	furnis
83	6580000	6000	1	0	1	0	0	1	1	0	0	0	1	0	0	0	1	
315	4095000	5600	0	0	0	0	0	0	1	0	0	0	1	0	0	0	1	
437	3290000	5880	1	0	0	0	0	0	1	0	0	0	0	1	0	0	0	
393	3500000	7424	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
247	4550000	8400	0	1	0	0	0	1	1	0	0	0	0	0	1	0	0	



Data split

```
In [ ]: from sklearn.linear_model import LinearRegression
        from sklearn.model_selection import train_test_split # Libreria para hacer split entre train y test
```

```

from sklearn.metrics import mean_squared_error

# Construccion de Los dataframes X e y para La regresion
# Variables independientes
X = df.drop(['price'],axis=1)
# Variables dependientes
y=df[['price']]

# Divide a Los dataframes en 70 y 30
X_train, X_test, y_train, y_test= train_test_split(X,y, test_size=0.3, random_state=1)

```

## El modelo de Regresion

```

In [ ]: # Ejecuta la regresion lineal

# Construye el regresor
regression_model=LinearRegression()

# Genera el fit de la regresion en con funcion fit
regression_model.fit(X_train,y_train)

```

```

Out[ ]: ▼ LinearRegression
LinearRegression()

```

```

In [ ]: # Obtiene el coeficiente de regresion de entrenamiento (R2)
regression_model.score(X_train,y_train)

```

```

Out[ ]: 0.6677872717804468

```

```

In [ ]: # Coeficiente de regresion de test
# Es el R2, o cuanto explica el modelo de la data de test
regression_model.score(X_test,y_test)

```

```

Out[ ]: 0.666227118877776

```

## Insights

- No se precibe una gran diferencia entre el R2 de este modelo de regresion y el del hecho durante la lecion

## Otras medidas del modelo

```

In [ ]: # Otras Medidas de Modelo
from sklearn.metrics import mean_absolute_error, mean_squared_error

y_pred= regression_model.predict(X_test)
# Suma absoluta de Las diferencias
mae = mean_absolute_error(y_test,y_pred)
# Suma absoluta de Las diferencias al cuadrado
mse = mean_squared_error(y_test, y_pred)
# Raiz cuadrada de la suma absoluta de Las diferencias al cuadrado
rmse = np.sqrt(mse)

```

```

In [ ]: print(f'Mean absolute error: {mae:.2f}')
print(f'Mean squared error: {mse:.2f}')
print(f'Root mean squared error: {rmse:.2f}')

```

Mean absolute error: 847499.75  
Mean squared error: 1377413021560.65  
Root mean squared error: 1173632.40

```
In [ ]: from sklearn import datasets, linear_model
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
from scipy import stats

X2 = sm.add_constant(X_train)
est = sm.OLS(y_train, X2)
est2 = est.fit()
print(est2.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	price	R-squared:	0.668			
Model:	OLS	Adj. R-squared:	0.652			
Method:	Least Squares	F-statistic:	42.80			
Date:	Thu, 30 Mar 2023	Prob (F-statistic):	8.53e-76			
Time:	16:49:47	Log-Likelihood:	-5801.5			
No. Observations:	380	AIC:	1.164e+04			
Df Residuals:	362	BIC:	1.171e+04			
Df Model:	17					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	1.578e+06	2.45e+05	6.436	0.000	1.1e+06	2.06e+06
area	254.2440	27.893	9.115	0.000	199.392	309.097
bedrooms_3	2.238e+05	1.54e+05	1.452	0.147	-7.94e+04	5.27e+05
bedrooms_4	2.669e+05	2.06e+05	1.299	0.195	-1.37e+05	6.71e+05
bathrooms_2	9.911e+05	1.39e+05	7.139	0.000	7.18e+05	1.26e+06
stories_2	3.762e+05	1.4e+05	2.681	0.008	1e+05	6.52e+05
stories_3	8.435e+05	2.58e+05	3.266	0.001	3.36e+05	1.35e+06
stories_4	1.575e+06	2.55e+05	6.176	0.000	1.07e+06	2.08e+06
mainroad_1	4.643e+05	1.78e+05	2.605	0.010	1.14e+05	8.15e+05
guestroom_1	2.657e+05	1.63e+05	1.630	0.104	-5.49e+04	5.86e+05
basement_1	4.011e+05	1.32e+05	3.036	0.003	1.41e+05	6.61e+05
hotwaterheating_1	1.108e+06	2.72e+05	4.078	0.000	5.74e+05	1.64e+06
airconditioning_1	6.868e+05	1.29e+05	5.333	0.000	4.34e+05	9.4e+05
parking_1	3.874e+05	1.4e+05	2.772	0.006	1.13e+05	6.62e+05
parking_2	5.528e+05	1.49e+05	3.704	0.000	2.59e+05	8.46e+05
prefarea_1	5.648e+05	1.43e+05	3.953	0.000	2.84e+05	8.46e+05
furnishingstatus_2	-1.344e+04	1.39e+05	-0.097	0.923	-2.87e+05	2.6e+05
furnishingstatus_3	-3.222e+05	1.54e+05	-2.094	0.037	-6.25e+05	-1.97e+04
=====						
Omnibus:	77.451	Durbin-Watson:	2.080			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	194.943			
Skew:	0.989	Prob(JB):	4.66e-43			
Kurtosis:	5.898	Cond. No.	3.49e+04			
=====						

Notes:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 3.49e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Insights

- La mayoría de las variables son estadísticamente significativas
- Al parecer hay un problema de multicolinearidad con las variables independientes.
- Una posible rta es que el numero de banos depende del numero de habitaciones que tiene el apartamento. A mayor numero de habitaciones, mayor numero de banos.

## importancia de variables

```
In [ ]: # Obtiene los coeficientes de la ecuacion
for idx, col_name in enumerate(X_train.columns):
    print('Coef {} = {}'.format(col_name, regression_model.coef_[0][idx]))
```

```
Coef area = 254.24400328624608
Coef bedrooms_3 = 223805.9089968321
Coef bedrooms_4 = 266940.17296603194
Coef bathrooms_2 = 991136.0265089186
Coef stories_2 = 376192.51984639396
Coef stories_3 = 843493.7551981165
Coef stories_4 = 1574539.669540347
Coef mainroad_1 = 464254.0462820121
Coef guestroom_1 = 265720.0627023259
Coef basement_1 = 401146.1781159133
Coef hotwaterheating_1 = 1108374.242582496
Coef airconditioning_1 = 686791.1215761018
Coef parking_1 = 387373.1306411324
Coef parking_2 = 552796.9471235962
Coef prefarea_1 = 564830.1558158637
Coef furnishingstatus_2 = -13442.77351442361
Coef furnishingstatus_3 = -322240.66035238776
```

```
In [ ]: # Obtiene los coeficientes en un Dataframe
# Se usa la funcion enumerate
# Tener en cuenta el uso de [0] ya que algunas variables vienen en arrays , para obtener el valor completo
coeff_data = pd.DataFrame()

#Intercept
coeff_data = coeff_data .append({'Coeff': regression_model.intercept_[0], 'Var': 'Intercept'}, ignore_index=True)

# Coeficientes
for idx, col_name in enumerate(X_train.columns):
    coeff_data = coeff_data.append({'Var': col_name, 'Coeff': regression_model.coef_[0][idx]}, ignore_index=True)
coeff_data
```

[illegible]

Out[ ]:

	Coeff	Var
0	1578427.468	Intercept
1	254.244	area
2	223805.909	bedrooms_3
3	266940.173	bedrooms_4
4	991136.027	bathrooms_2
5	376192.520	stories_2
6	843493.755	stories_3
7	1574539.670	stories_4
8	464254.046	mainroad_1
9	265720.063	guestroom_1
10	401146.178	basement_1
11	1108374.243	hotwaterheating_1
12	686791.122	airconditioning_1
13	387373.131	parking_1
14	552796.947	parking_2
15	564830.156	prefarea_1
16	-13442.774	furnishingstatus_2
17	-322240.660	furnishingstatus_3

Insights

- En promedio, un aumento de 1m cuadrado genera un aumento de \$254 en el precio de la vivienda.
- Los coeficientes que representan una mayor impacto al precio de la vivienda en promedio son:
  - Tener agua caliente y,
  - Que la vivienda este ubicada en un cuarto piso
- Como era de esperarse, el hecho de que la vivienda no este o parcialmente amueblada indica una reduccion en el precio de la vivienda.

predicción de los resultados del modelo

In [ ]:

```
# Comparacion de Prediccion vs Actual

df_preds = pd.DataFrame({'Actual':y_test.squeeze(), 'Predicted':y_pred.squeeze()})
print(df_preds)
```


	Actual	Predicted
244	4550000	5213511.664
132	5810000	5105398.344
385	3570000	2645889.026
133	5810000	4694902.856
481	2940000	2264886.277
..	...	...
330	3990000	4358017.984
542	1750000	2640804.146
433	3290000	3480259.650
80	6629000	6215278.719
46	7525000	7431791.391

[163 rows x 2 columns]

### Visualización de los resultados del modelo de regresión

```
In [ ]: plt.scatter(y_test, y_pred)
plt.title('Comparacion de Prediccion (x) vs Real (y)')
```

```
Out[ ]: Text(0.5, 1.0, 'Comparacion de Prediccion (x) vs Real (y)')
```



### Bonus: Regression with standardized coefficients

```
In [ ]: from sklearn.preprocessing import StandardScaler
df1=df.copy()
columns_df1= df1.columns.to_list()

scaler = StandardScaler()
for column in columns_df1:
    df1[column]=StandardScaler().fit_transform(df1[[column]])

df1.sample(5)
```



Out [ ]:

	price	area	bedrooms_3	bedrooms_4	bathrooms_2	stories_2	stories_3	stories_4	mainroad_1	guestroom_1	basement_1	hotwaterheating_1	airconditioning_1	parking_1	parking_2	prefarea_1	furnishingstatus_2	furnish
520	-1.244	1.172	-1.111	-0.495	-0.601	-0.883	-0.278	-0.286	0.403	-0.466	-0.737	-0.220	-0.681	-0.550	-0.533	-0.555	-0.848	
243	-0.120	-1.201	0.900	-0.495	-0.601	1.132	-0.278	-0.286	0.403	-0.466	1.358	-0.220	-0.681	-0.550	-0.533	-0.555	-0.848	
432	-0.795	0.417	0.900	-0.495	-0.601	-0.883	-0.278	-0.286	0.403	2.144	1.358	-0.220	-0.681	-0.550	-0.533	-0.555	-0.848	
401	-0.682	2.002	0.900	-0.495	-0.601	1.132	-0.278	-0.286	0.403	-0.466	-0.737	-0.220	-0.681	-0.550	1.877	1.801	-0.848	
225	-0.011	0.592	-1.111	-0.495	-0.601	-0.883	-0.278	-0.286	0.403	-0.466	-0.737	-0.220	1.469	-0.550	1.877	-0.555	1.180	

In [ ]:

```
# Construccion de los dataframes X e y para la regresion
# Variables independientes
X1 = df1.drop(['price'],axis=1)
# Variables dependientes
y1=df1[['price']]

# Divide a los dataframes en 70 y 30
X1_train, X1_test, y1_train, y1_test= train_test_split(X1,y1, test_size=0.3, random_state=1)
```

In [ ]:

```
# Ejecuta la regresion lineal
# Construye el regresor
regression_model=LinearRegression()

# Genera el fit de la regresion en con funcion fit
regression_model.fit(X1_train,y1_train)
```

Out [ ]:

LinearRegression

LinearRegression()

In [ ]:

```
# Obtiene el coeficiente de regresion de entrenamiento (R2)
regression_model.score(X1_train,y1_train)
```

Out [ ]:

0.667787271780447

In [ ]:

```
# Coeficiente de regresion de test
# Es el R2, o cuanto explica el modelo de la data de test
regression_model.score(X1_test,y1_test)
```

Out [ ]:

0.6662271188777761

In [ ]:

```
X2 = sm.add_constant(X1_train)
est = sm.OLS(y1_train, X2)
est2 = est.fit()
print(est2.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.668
Model:                  OLS      Adj. R-squared:             0.652
Method:                 Least Squares      F-statistic:         42.80
Date:                   Thu, 30 Mar 2023    Prob (F-statistic):    8.53e-76
Time:                   16:50:50    Log-Likelihood:       -314.25
No. Observations:       380      AIC:                   664.5
Df Residuals:           362      BIC:                   735.4
Df Model:               17
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                -0.0057       0.029      -0.194      0.846      -0.063      0.052
area                  0.2954       0.032       9.115      0.000       0.232      0.359
bedrooms_3            0.0596       0.041       1.452      0.147      -0.021      0.140
bedrooms_4            0.0568       0.044       1.299      0.195      -0.029      0.143
bathrooms_2           0.2342       0.033       7.139      0.000       0.170      0.299
stories_2              0.0999       0.037       2.681      0.008       0.027      0.173
stories_3              0.1166       0.036       3.266      0.001       0.046      0.187
stories_4              0.2227       0.036       6.176      0.000       0.152      0.294
mainroad_1             0.0862       0.033       2.605      0.010       0.021      0.151
guestroom_1            0.0545       0.033       1.630      0.104      -0.011      0.120
basement_1             0.1026       0.034       3.036      0.003       0.036      0.169
hotwaterheating_1      0.1244       0.030       4.078      0.000       0.064      0.184
airconditioning_1      0.1711       0.032       5.333      0.000       0.108      0.234
parking_1              0.0875       0.032       2.772      0.006       0.025      0.150
parking_2              0.1228       0.033       3.704      0.000       0.058      0.188
prefarea_1             0.1284       0.032       3.953      0.000       0.065      0.192
furnishingstatus_2     -0.0035       0.037      -0.097      0.923      -0.076      0.069
furnishingstatus_3     -0.0809       0.039      -2.094      0.037      -0.157     -0.005
=====
Omnibus:               77.451    Durbin-Watson:          2.080
Prob(Omnibus):          0.000    Jarque-Bera (JB):       194.943
Skew:                   0.989    Prob(JB):               4.66e-43
Kurtosis:               5.898    Cond. No.                3.33
=====
```

Notes:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [ ]:
```